

# Foundations of Distributed Systems

José Orlando Pereira

HASLab / Departamento de Informática  
Universidade do Minho



2018/2019

# Motivation

- Handle a large number of clients and requests with a single server
- The “c10k problem” in 1999:
  - <http://www.kegel.com/c10k.html>
- Examples:
  - financial, games, ...
  - notifications in mobile apps
  - machine-to-machine (M2M)

# Case study

- Simple chat server:
  - Forward all messages to all clients
- Consider:
  - Large number of clients
  - Slow connections



# First threaded solution

- For each connection:
  - Handler thread
- When reading, write to all other connections
- Use buffering:
  - To minimize system calls
  - To cope with slow readers

# Sockets in java.net

```
ServerSocket ss=new ServerSocket(12345);

while(true) {
    Socket s=ss.accept();

    InputStream is=s.getInputStream();
    OutputStream os=s.getOutputStream();

    // i/o

    s.close();
}
```

# Buffers in java.net

```
ServerSocket ss=new ServerSocket(12345);

while(true) {
    Socket s=ss.accept();

    InputStream is=new BufferedInputStream(s.getInputStream());
    OutputStream os=new BufferedOutputStream(s.getOutputStream());

    // i/o

    os.flush();
    s.close();
}
```

Needed to  
actually write!

# Memory

- Memory used:  $n$  connections x messages in transit ( $\sim n^2$ )
  - Caused by data copying in stacked abstractions
    - Serialization!
  - Overhead in allocation and garbage collection
- Solution: Store transient data in reusable shared buffers
  - Pointers/indexes into statically allocated data


# Sockets in java.nio

```
ServerSocketChannel ss=ServerSocketChannel.open();  
ss.bind(new InetSocketAddress(12345));  
  
while(true) {  
    SocketChannel s=ss.accept();  
  
    // i/o  
  
    s.close();  
}
```




# Buffers in java.nio

- Buffer = Array + Indexes:




# Buffers in java.nio

- Put/read: advances position, sets content




# Buffers in java.nio

- Get/write: advances position, gets content

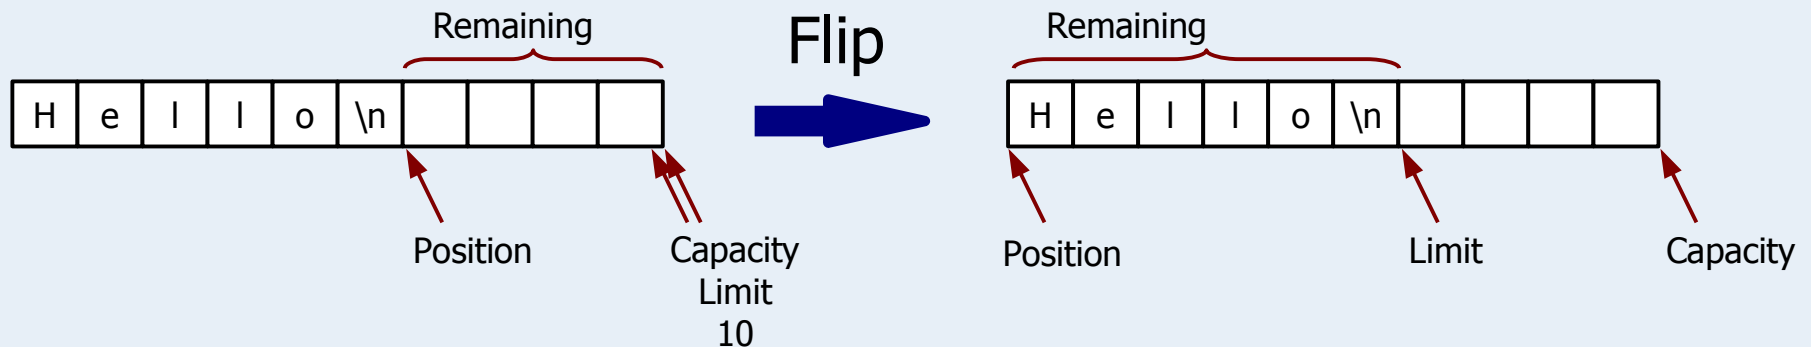



# Buffers in java.nio



- Duplicate and wrap: multiple pointers into the same array



# Buffers in java.nio



# Buffers in java.nio



# Shared buffers

- Memory used: messages in transit ( $\sim n$ )
- Ideally, never allocate or dispose of memory in normal operation:
  - No overhead, but...
  - Needs reference counting to know when to reuse

# Flushing buffers

```
ServerSocket ss=new ServerSocket(12345);

while(true) {
    Socket s=ss.accept();

    InputStream is=new BufferedInputStream(s.getInputStream());
    OutputStream os=new BufferedOutputStream(s.getOutputStream());

    // i/o

    os.flush();
    s.close();
}
```

What if  
write blocks?



# Second threaded solution

- For each connection:
  - Reader thread
  - Writer thread
  - Pending queue
- When reading, insert in queues and notify writer threads
- When writing, remove from queue and notify reader threads

# Threads

- Problems:
  - Memory overhead (stacks)
  - Context switches and lock contention
  - “Thundering herd”, hidden queue, and fairness

# Blocking sockets

```
try {  
    ByteBuffer buf=ByteBuffer.allocate(100);  
  
    s.read(buf);  
    buf.flip();  
  
    r.write(buf);  
} catch(IOException e) {  
    report(e);  
}
```

# Asynchronous sockets

```
ByteBuffer buf=ByteBuffer.allocate(100);

s.read(buf, null, new CompletionHandler() {
    public void completed(Integer result, Object a) {
        buf.flip();

        r.write(buf, ...);
    }
    public void failure(Throwable t, Object a) {
        report(t);
    }
});
```

# Thread pools

- For non-blocking, short-lived events:
  - One pool thread for hardware thread
- While all threads are blocked, the application stops handling events

```
AsynchronousChannelGroup g =  
    AsynchronousChannelGroup.withFixedThreadPool(...);  
  
AsynchronousSocketChannel s =  
    AsynchronousSocketChannel.open(g);  
  
... /* callbacks use g.shutdown() to exit */  
  
g.awaitTermination(Long.MAX_VALUE, TimeUnit.SECONDS);
```