


Crash and Recovery

- Asynchronous system model
- Crash and recovery events (a.k.a reboot)
 - Assume that there is always recovery
- Two types of variables:
 - Volatile, lost on crash
 - Persistent, kept after recovery

Operations

- The application changes volatile variables:




Operations

- Changes can now be copied to persistent variables:

Volatile:




Persistent:



Failure and restart

- Restart erases volatile variables:

Volatile:




Persistent:



Failure and restart

- Volatile variables are restored from the persistent copy:



Challenges

- Single process failure and recovery
- Partial system failure and recovery

Redo log

- Assume a persistent sequential log:


Volatile variables:

	1	2	3			6	...								
--	---	---	---	--	--	---	-----	--	--	--	--	--	--	--	--

Persistent variables:

0	1	2	3	...											
---	---	---	---	-----	--	--	--	--	--	--	--	--	--	--	--


Sequential log:



Redo log

- Changed variables are copied to the log followed by a commit marker:


Volatile:



Persistent:



Log:




Redo log

- Changes can now be copied to persistent variables:

Volatile:



Persistent:



Log:



Redo log

- Eventually all persistent variables are updated:

Volatile:



Persistent:



Log:



Redo log

- In case of restart, upon recovery, redo all changes from the log:



Volatile:



Persistent:



Log:



Redo log

- And restore volatile variables to become operational again:



Volatile:



Persistent:



Log:

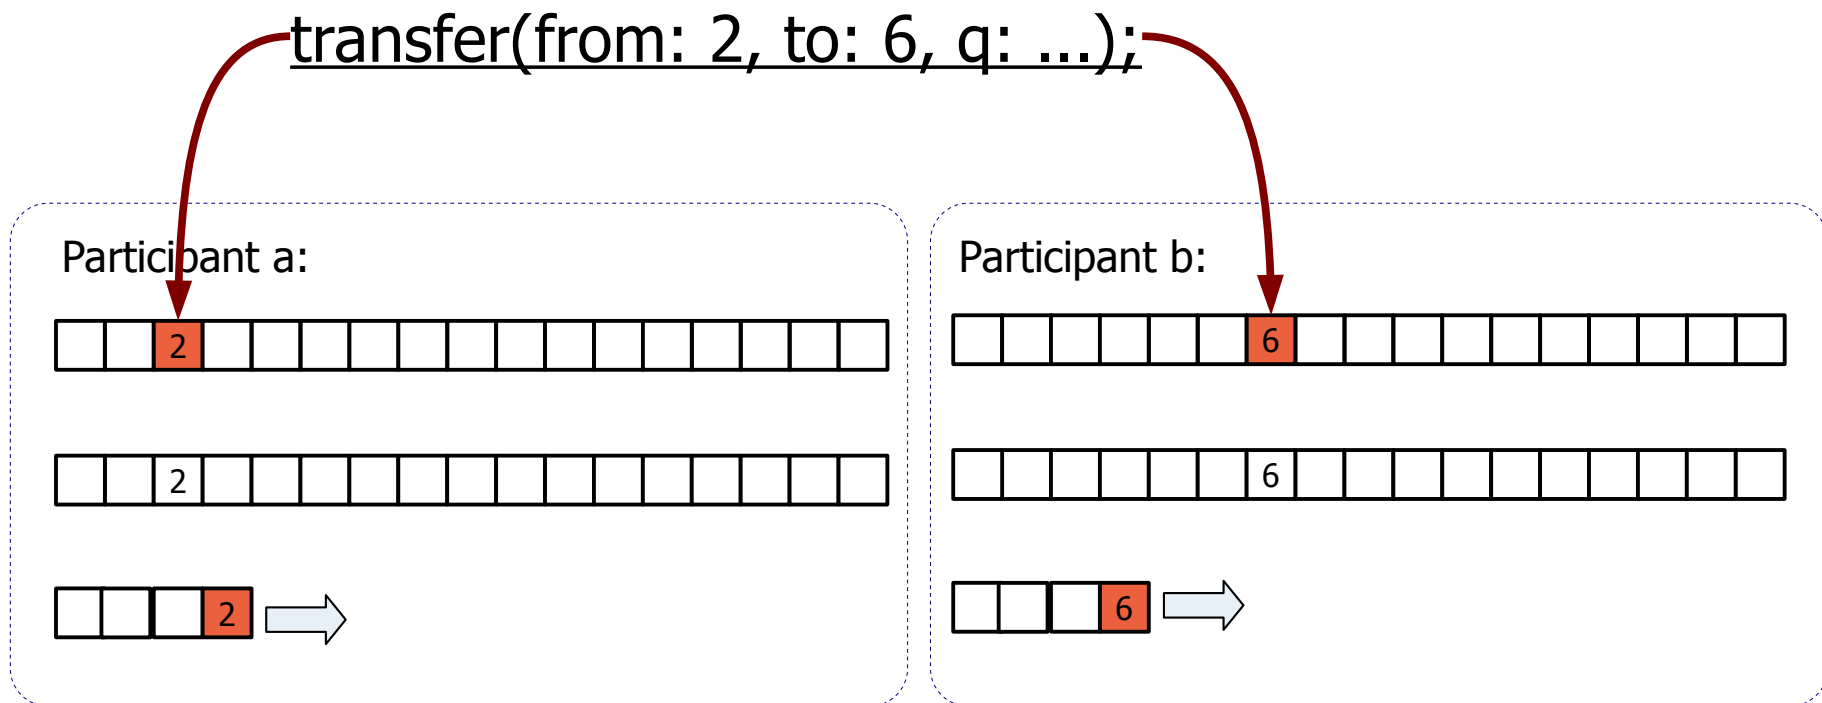


Challenges

- ~~Single process failure and recovery~~
- Partial system failure and recovery

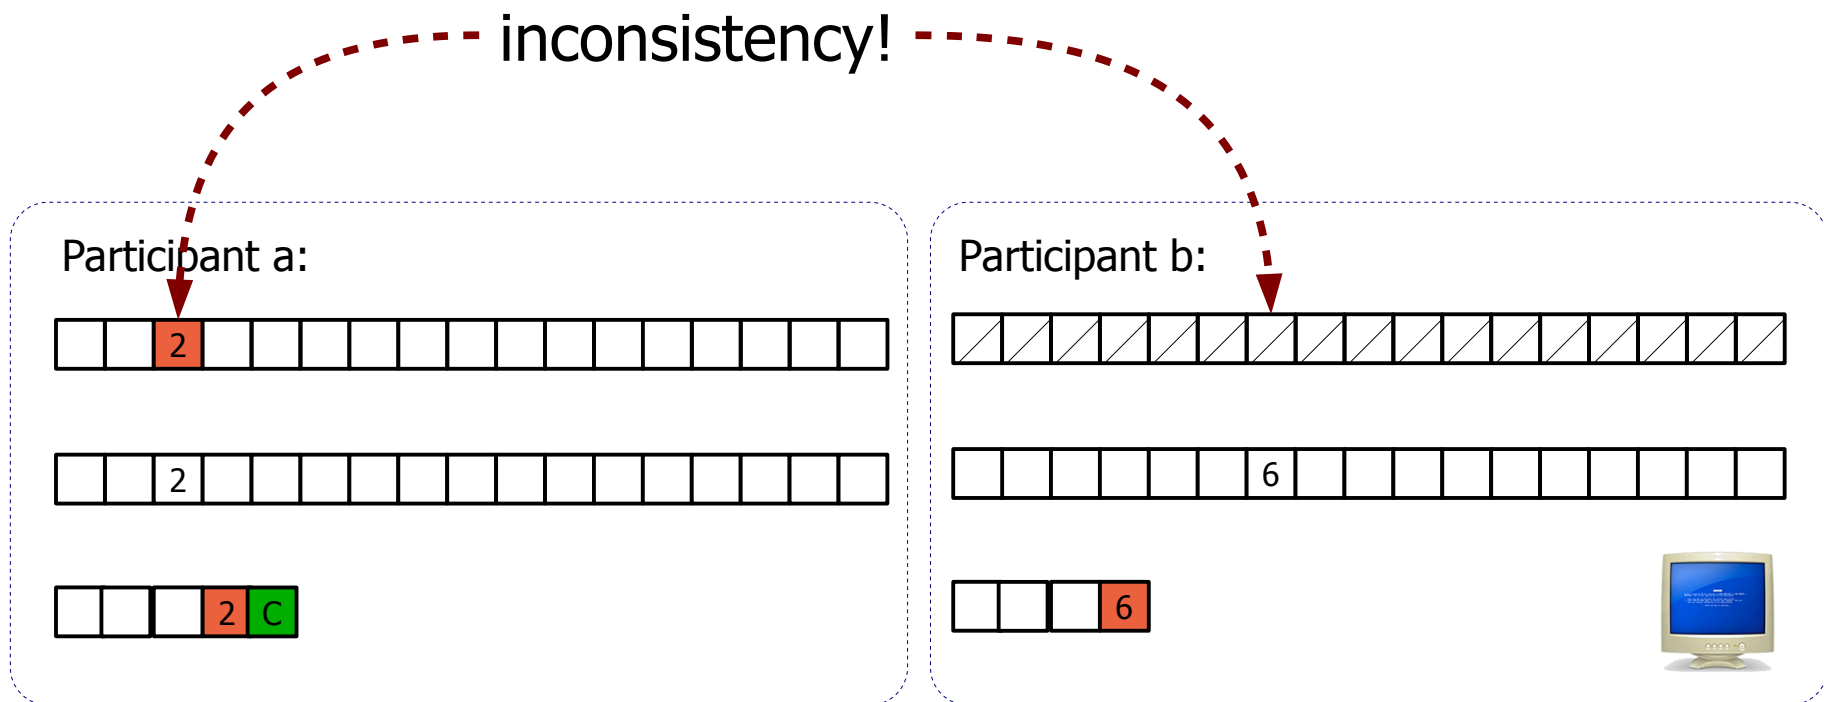
Distributed operations

- Distributed operation updates variables in multiple processes:



Partial failure and restart

- If each participant commits independently, then...



Two Phase Commit (2PC)

- Add a coordinator with an additional persistent log:

Coordinator log:



Participant a:



Participant b:




2PC: Phase 1

- The coordinator asks all participants to prepare, inserting prepared markers in logs:


Coordinator log:



Participant a:



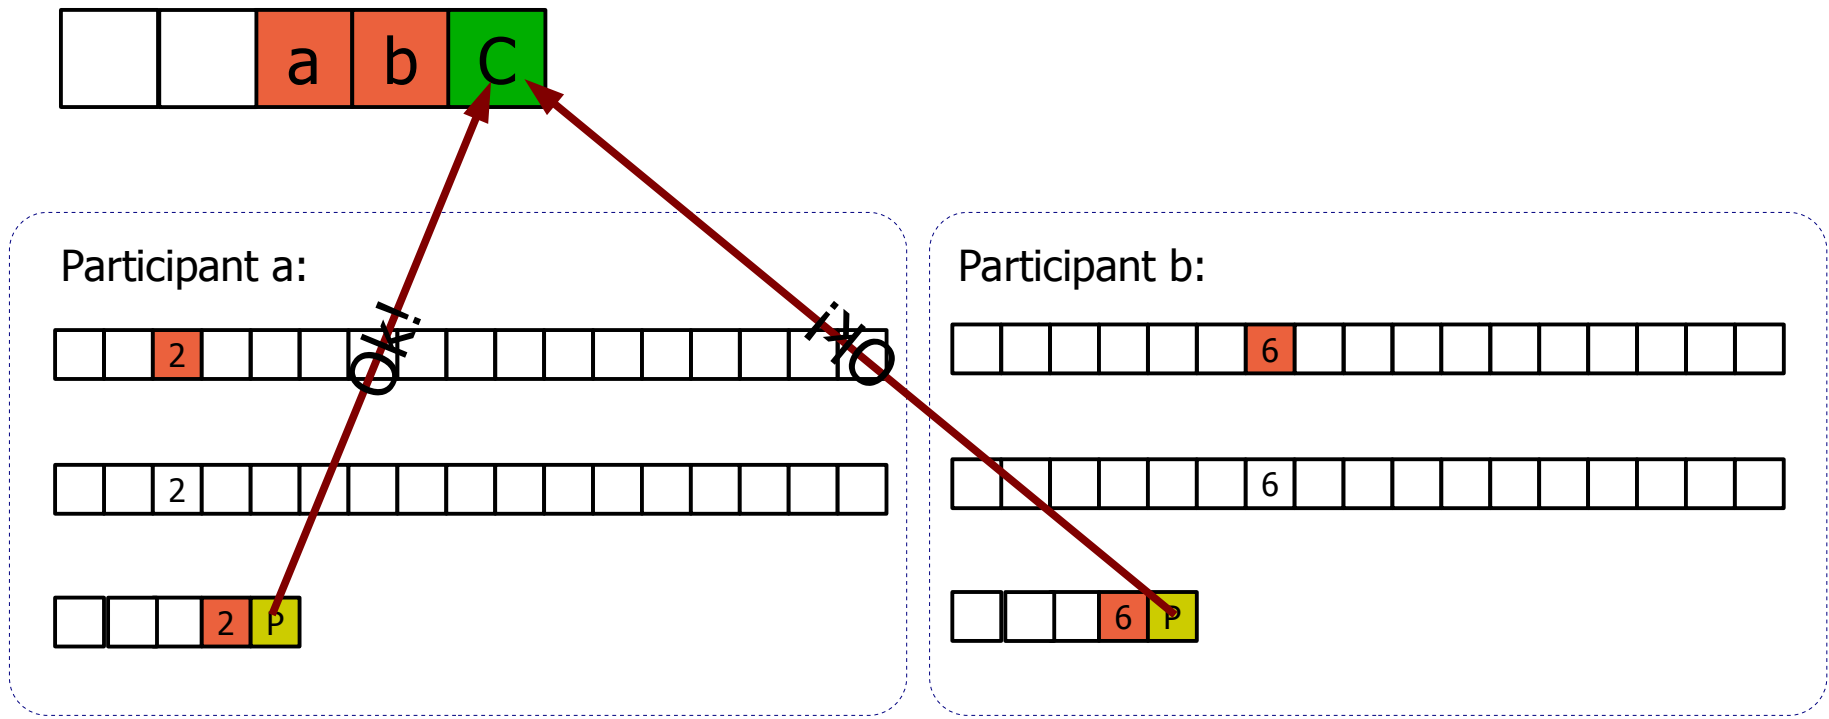
Participant b:



2PC: Phase 1 Completed

- If all participants are able to prepare, the coordinator inserts a commit marker:

Coordinator log:




2PC: Phase 2

- Participants are informed of outcome inserting commit markers in their logs:

Coordinator log:



Participant a:



Participant b:



Commit

Commit


2PC: Phase 1 Aborted

- If a single participant calls for abort or fails the entire transaction is marked for rollback:


Coordinator log:



Participant a:



Participant b:




2PC: Phase 2 Aborted

- If a single participant calls for abort or fails the entire transaction is marked for rollback:


Coordinator log:



Participant a:



Participant b:



Rollback

2PC: Recovery

- Restart of the coordinator:
 - Before 2PC started: Abort, as it might have missed a resource
 - During 2PC: Restart current phase by repeating the request
- Restart of a participant:
 - Has not voted: Local rollback, will abort the entire transaction
 - Has voted: Wait for the decision from the coordinator

Summary

