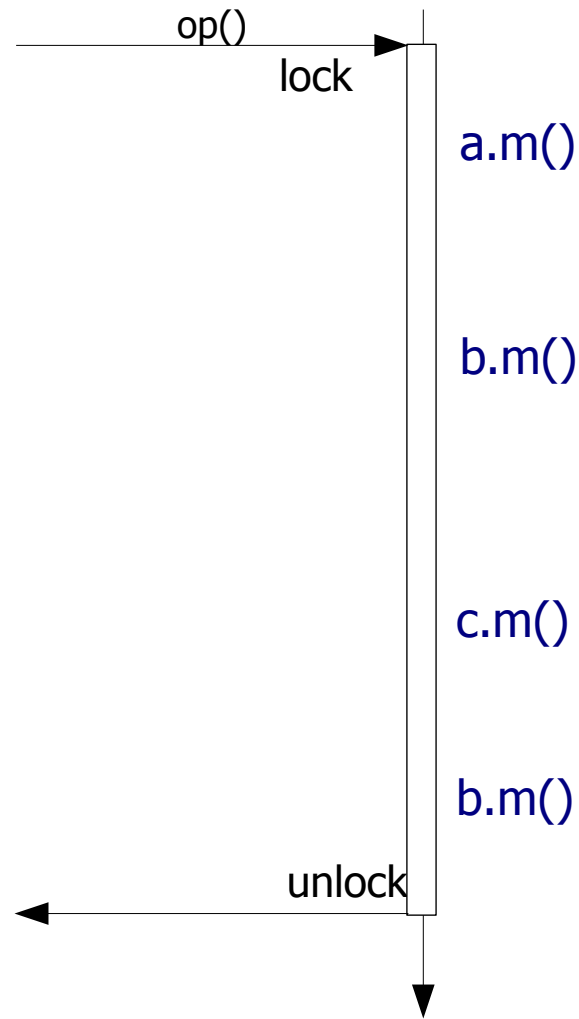


Challenges


- Atomic operations on multiple distributed objects
 - Concurrent client invocations
 - Client and server failure

Locking



- Lock/unlock operations in a lock server
- A single global lock protects all data items
- Acquired for the duration of the composite operation

Locking




- Operations never overlap
- History is a sequence of operations

Summary


- Performance bottleneck:
 - At most one operation executing
- Remote invocations needed for lock / unlock
- Single point of failure:
 - Lock server fails
 - Lock client fails without unlock()

Locking



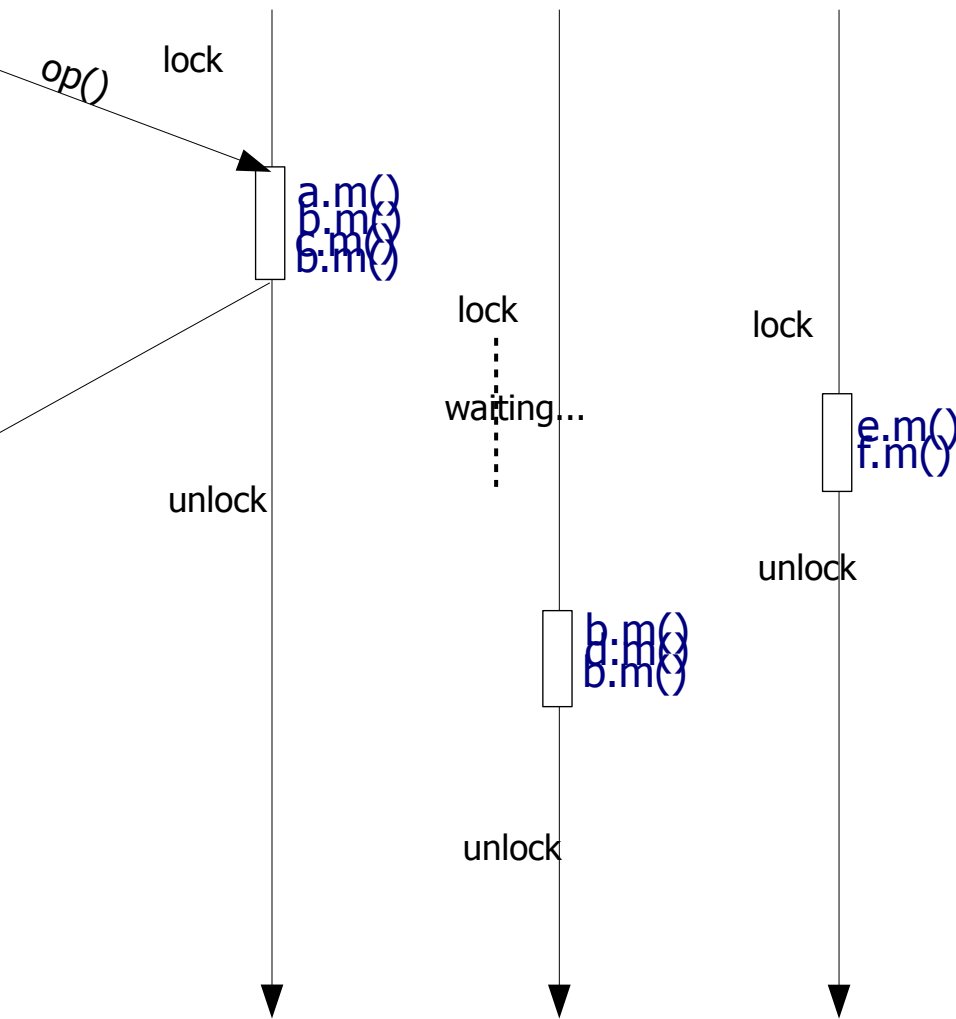
- Multiple lock objects
- Each data item is protected by a lock
- All locks are acquired during the operation

Locking



- Operations on different items may overlap

Locking




- Indistinguishable from:
 - Instantaneous execution with delay on round-trip
- Clients cannot deny that operations do not overlap

Summary

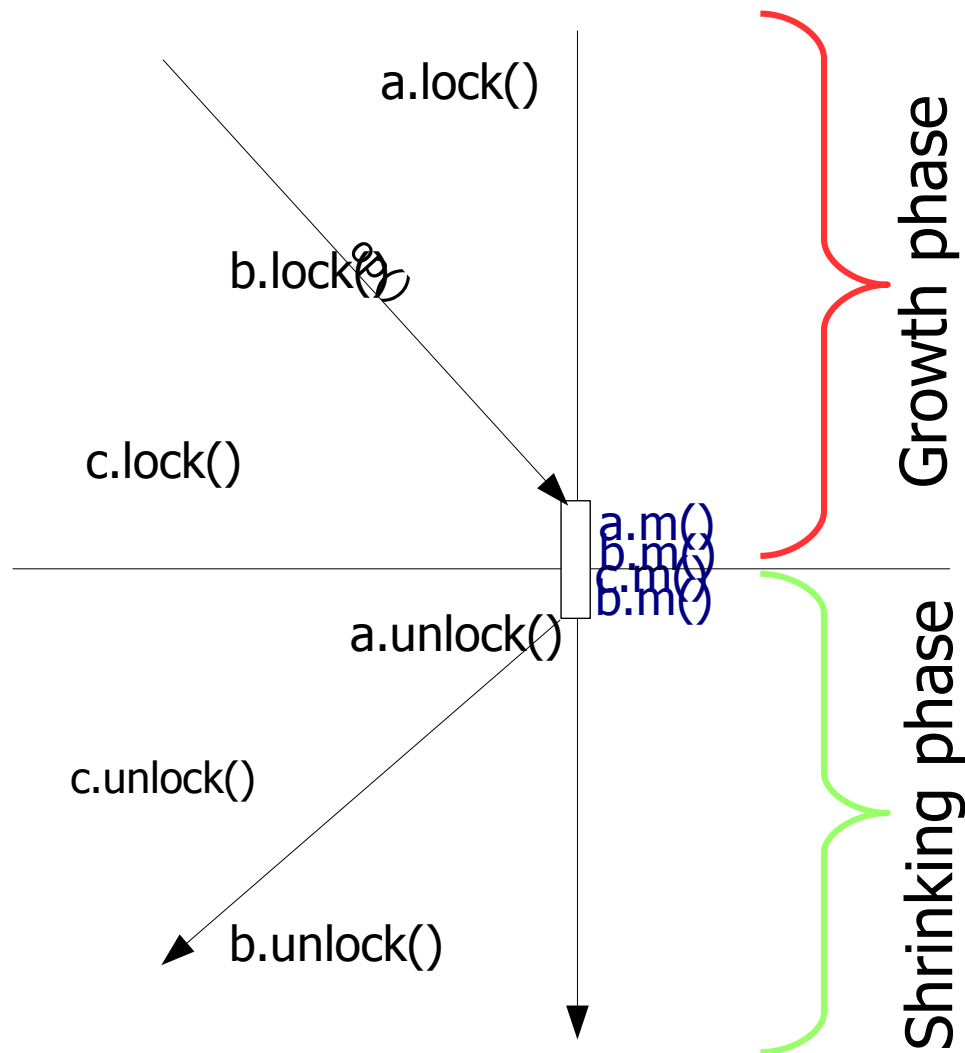
- ~~Performance bottleneck~~
- Remote invocations needed for lock / unlock
- Single point of failure:
 - ~~Lock server fails~~
 - Lock client fails without unlock()
- Must know all items accessed on start
- Deadlocks

2-Phase locking



- All lock requests precede all unlock requests
- Lock can be implicit in operation

2-Phase locking

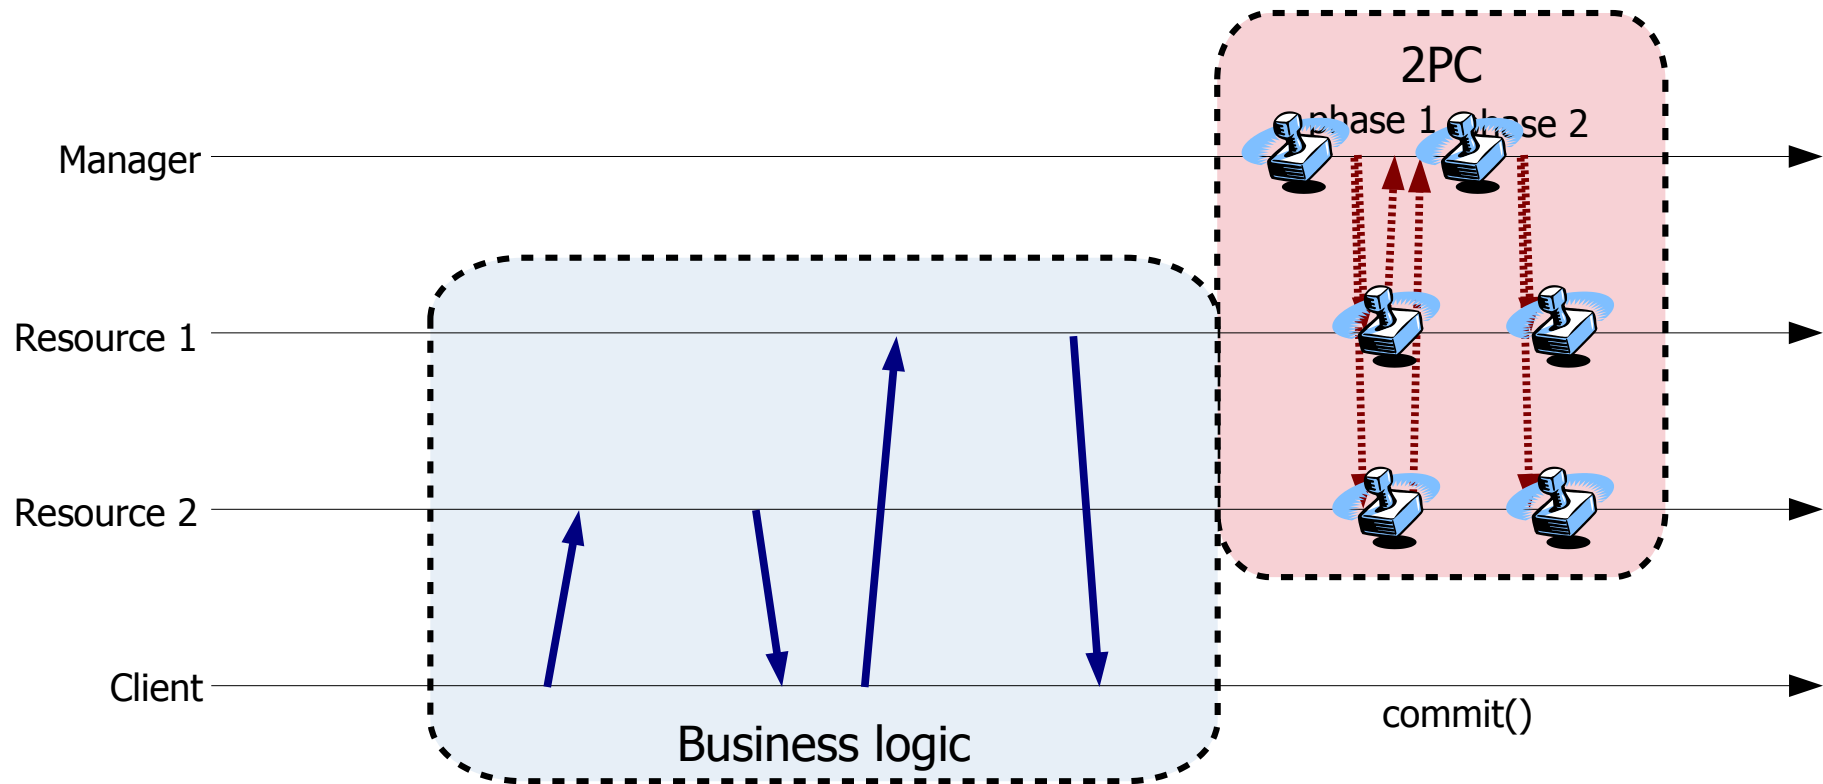


- An observer cannot deny that all happens in-between phases

Summary

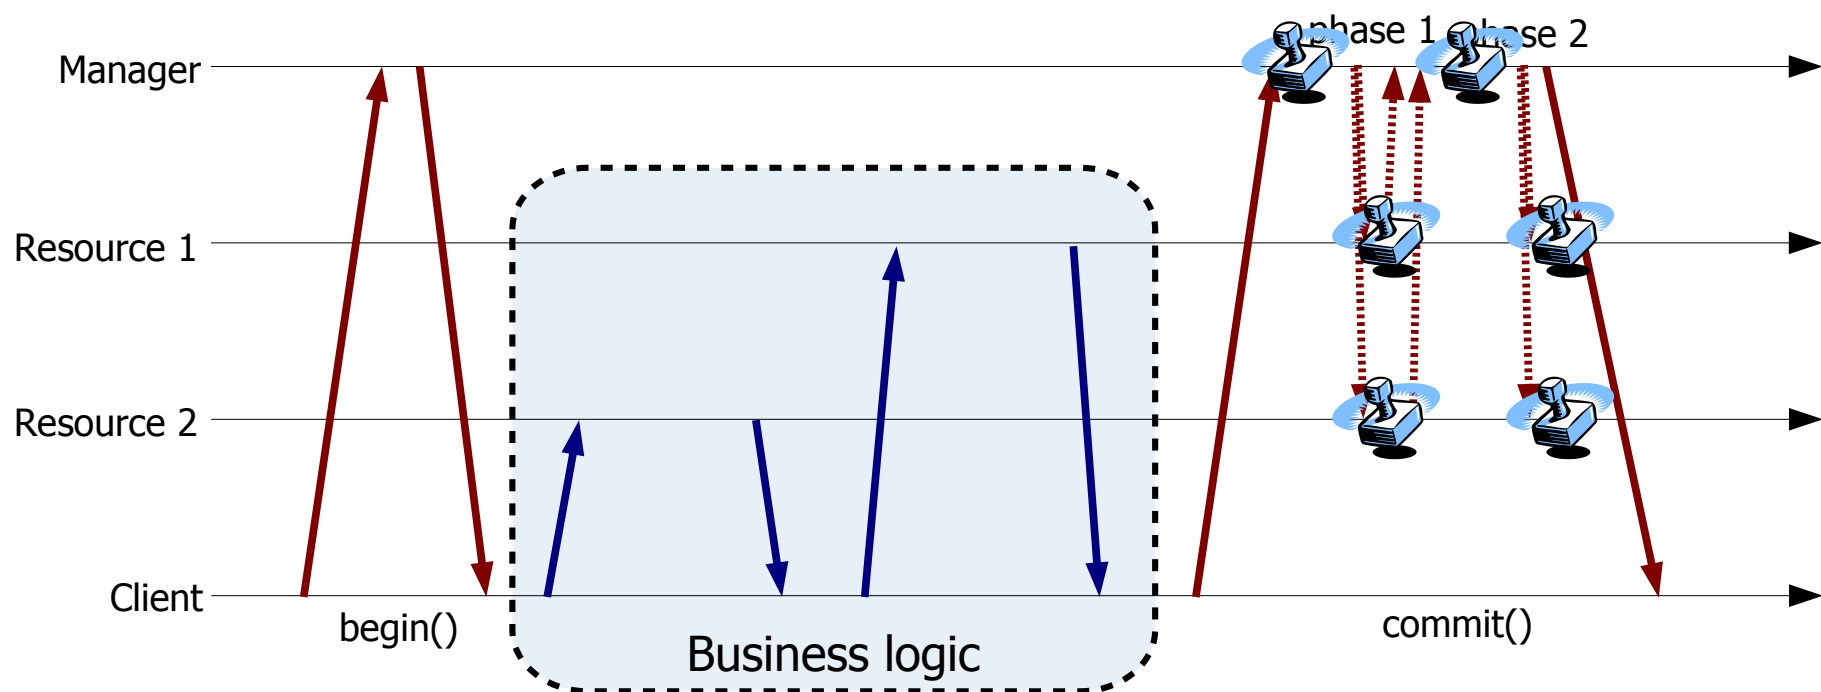
- ~~Performance bottleneck~~
- Remote invocations needed for ~~lock~~ / unlock
 - Lock operation is implicit
- Single point of failure:
 - ~~Lock server fails~~
 - Client fails without unlock()
- ~~Must know all items accessed on start~~
- Deadlocks

Application vs Transactions




- How to connect them?
 - Don't change business logic

Transaction demarcation




- Call transaction manager before/after
- How does manager discover participants?


Transactional RPC + 2PC



Transactional RPC + 2PC



Transactional RPC + 2PC + 2PL




Summary

- ~~Performance bottleneck~~
- ~~Remote invocations needed for lock / unlock~~
 - Lock operation is implicit
 - Unlock operation is implicit (txn. commit)
- Single point of failure:
 - ~~Lock server fails~~
 - ~~Lock client fails without unlock()~~ (rollback)
- ~~Must know all items accessed on start~~
- ~~Deadlocks (rollback)~~

Summary




- Atomicity with faults: 2PC
- Atomicity with concurrent clients: 2PL
- With 2PL + 2PC:
 - Rollback on deadlock / client failure
 - Implicitly release locks on commit / rollback

2PC+2PL: Challenges




- Scalability with:
 - More resources?
 - Slower 2PC (remote resources)?

Store and Forward




Transactional queues



- Insert and remove from temporary DBs abstracted as transactional send/receive
- Advantages:
 - Performance and fault decoupling
 - Asynchronous (pull) operations


Server cluster

- Multiple servers polling the same queue:
 - Each message processed exactly once
 - No order enforced



Request/Reply


- Client/server:
 - Server polls for requests
 - Client polls for replies
- Multiple transaction contexts:



Limitations of queues

- Consider adding another application...
- Addressing is explicit in applications:
 - Each sender must know exactly which receivers are interested in a message

Message brokers



- Goal: Decouple message routing from sending and receiving
- Receivers subscribe to messages that fit some criteria
- Senders publish messages

Topic-based routing

- Publisher tags messages with a topic id
- Subscriber requests a topic
- Options:
 - Flat topics
 - Hierarchical topics

Content-based routing

- Publisher structures message content using a common structure
- Subscriber registers a content filter
- Options:
 - Relational data / SQL WHERE clauses
 - XML / XQuery

Summary

- Transactional messaging decouples:
 - Faults
 - Performance
 - Application logic
- Two options:
 - Tightly coupled applications: 2PL + 2PC + TRPC
 - Loosely coupled: 2PC + MQ